

# Enterprise **vs** Non-Enterprise AI Agents

# Intro

- Enterprise AI Agents
- Non-Enterprise AI Agents

# What Are AI Agents? Why Now?

- AI Agents
- Why now?
- Analogy

# Enterprise AI Agent Frameworks

Framework	Purpose	Fun Fact
IBM Watsonx.ai + Governance	AI compliance, orchestration, and explainability	Used by <b>80%</b> of <b>Fortune 100</b>
Azure AI Agents + Copilot Stack	Enterprise workflows, embedded in M365	Powers over <b>300M</b> users
AWS Agents on Bedrock / SageMaker	Build scalable agent workflows	Bedrock hit <b>10K+</b> customers in <b>6 months</b>
Google Vertex AI + LangChain (Ent)	Real-time agents with grounded outputs	Runs on <b>TPUv5</b> , trillions of params
K8s Agent Systems	Enterprise CI/CD and API-based orchestration	Kubernetes powers <b>96%</b> of enterprise containers

# Non-Enterprise AI Agent Frameworks

Framework	Use	Fun Fact
<b>Auto-GPT / BabyAGI / CAMEL</b>	Loop-based, goal-driven automation	Auto-GPT repo hit 120K stars in 2 months
<b>LangChain OSS</b>	Lightweight agent chaining	70% of Gen-AI hackathon winners used it
<b>CrewAI / Agent-LLM</b>	Multi-agent teams with memory	CrewAI agents can cook meals
<b>ReAct + LlamaIndex</b>	Reason + Action loop + vector search	Inspired by human psychology
<b>Meta's OpenAGI / HuggingGPT</b>	Connect 100+ models to act as one	HuggingGPT can 'hug' over 100 tools

# Enterprise **vs** Non-Enterprise: The Real Differences

Feature	Enterprise	Non-Enterprise
Security	Strict RBAC, audit, logs	Minimal or none
Model Access	Fine-tuned + private LLMs	Public APIs and OSS
Orchestration	CI/CD, hybrid cloud	Localhost, Colab
Governance	Explainability dashboards	Not a focus
Team Use	Cross-org, monitored	Single-user or team testing

# Architecture Tips

Whether you're an enterprise architect or a weekend hacker, here's how to approach AI agents effectively

## **Always define**

- Your Agent's Role
- Tools it can access
- Memory, it needs
- Feedback loop (success/failure)

## **For enterprises:**

- Include AI governance like Watsonx.governance
- Use containerized orchestration (OpenShift / EKS)
- Monitor performance, data lineage

## **For builders:**

- Use LangChain + LlamaIndex + CrewAI combo
- Set clear prompt boundaries
- Use local models where possible for cost control

# Thank You!



Framework	Purpose	Fun Fact
<b>IBM Watsonx.ai + Watsonx.governance</b>	Enterprise LLMs, data lineage, model risk	Used by 80% of Fortune 100
<b>Azure AI Agents + M365 Copilot</b>	Seamless AI orchestration across Teams & Office	300M+ users interact with it
<b>AWS Agents on Bedrock/SageMaker</b>	Multi-agent pipelines with guardrails	10K+ customers in < 6 months
<b>Google Vertex AI + LangChain (Enterprise)</b>	Grounded Gen-AI outputs, tool use	Runs on TPUv5 with trillions of params
<b>Kubernetes-based Open Frameworks</b>	Containerized, observable AI agent systems	96% of enterprises run on Kubernetes